

Langchain - Detailed Explanation

what is LangChain?

- **Langchain** is a **framework** used to build applications powered by **Large Language Models**
- LangChain is a powerful framework designed to build applications using Large Language Models (LLMs).
- It helps developers connect language models with external data sources, tools, memory, and logic in a structured and reusable way.
- Traditional LLM usage is limited to prompt-response interactions, but LangChain extends this by enabling document understanding, retrieval-based question answering, conversational memory, tool usage, and agent-based reasoning.
- In real-world applications, data is often stored in PDFs, databases, websites, or APIs.
- LangChain acts as a bridge between raw data and intelligent AI responses. It is widely used for building chatbots, document Q&A systems, search engines, summarizers, and AI assistants.
- like GPT,Gemini,Claude,etc
- it helps in:
 - Handling Long documents
 - Connecting models+data+pr`mompts+tools
 - Building chatbots,RAG systems,agents,assistants

Load the Documents-Documents Loaders

why needed?

- LLMs cannot directly read files like PDF,Word,CSV,websites,etc.
- so Langchain provides Document Loaders
- Large Language Models like GPT, Gemini, or LLaMA are powerful but stateless and context-limited.
- They cannot:
 - Remember long conversations
 - Access private documents directly
 - Perform multi-step reasoning reliably
 - Interact with tools or APIs by default

what they do:

- Load raw data From different sources
- Convert it into a standard document format

Document Loaders (In-Depth)

- Document Loaders are responsible for loading data from different sources and converting it into text format that LLMs can understand.
- Supported sources include:
 - PDF files
 - Word documents
 - Text files
 - CSV and Excel files
 - Websites and APIs
 - Databases

Examples

- PDF
- TXT
- DOCX
- CSV
- Web URLs
- Databases

Common Loaders:

- PyPDFLoader
- TextLoader
- CSVLoader
- WebBaseLoader

LangChain solves these problems by:

- Adding memory to conversations
- Connecting LLMs to vector databases
- Structuring prompts systematically
- Enabling reasoning through chains and agents

This makes LangChain suitable for enterprise-grade AI applications.

Langchain Load Different Models(LLMs)

Large Language Models (LLMs)

- LLMs generate the final response using the prompt and retrieved context.
- LLMs do not search data themselves; they rely entirely on provided context, making retrieval accuracy critical.

why need

- different applications need different models:
 - Chatbots
 - Summarization
 - Q&A
 - Code generation

Models LangChain Supports:

- OpenAI(GPT-4, GPT-3.5)
- Google Gemini
- Anthropic Claude
- HuggingFace models
- Local LLMs (LLaMA, Mistral)

Role:

- Generate text
- Answer questions
- Reason over data

Langchain provides a common interface, so model switching is easy

Provide Structure – Prompt Templates

Problem without prompt templates:

- Hard-Coded prompts are:
 - Messy
 - Not reusable
 - Error-prone

What Prompt Templates do:

- Create structured prompts
- Insert variables dynamically
- Make prompts reusable

Benefits

- Clean code
- Dynamic inputs

- Easy debugging

Divide Data into Multiple Parts – Chunking

Chunking (Text Splitting)

- Chunking is the process of dividing long text into smaller, overlapping chunks.
- Key concepts:
 - Chunk size: Number of characters or tokens
 - Chunk overlap: Shared content between chunks to maintain context
- Common text splitters:
 - RecursiveCharacterTextSplitter
 - CharacterTextSplitter
 - TokenTextSplitter

Why chunking is important?

- LLMs have token limits
- Long documents cannot be passed directly

Chunking does:

- Splits large documents into small pieces
- Keeps overlap for context continuity
- Proper chunking improves retrieval accuracy and ensures important context is not lost.

Example

- Chunk size: 500 tokens
- Overlap: 50 tokens

Result

- Document → chunks → embeddings **this is mandatory for RAG systems**

Store Embeddings – Vector Database

What are embeddings?

- Embeddings = numerical representation of text meaning
- Embeddings convert text into numerical vectors that represent semantic meaning.
- Similar texts produce similar vectors.

Why embeddings are important:

- Enable similarity search
- Allow semantic understanding instead of keyword matching

Popular embedding models:

- OpenAI Embeddings
- HuggingFace Sentence Transformers
- Google Gemini Embeddings
- Each chunk is converted into an embedding vector and stored for future retrieval.

Why store embeddings?

- Fast similarity search
- Semantic retrieval
- Context-aware answers

Vector Databases:

- Vector databases store embeddings efficiently and allow fast similarity search.
- Common vector databases: - FAISS - Chroma - Pinecone - Weaviate - Milvus

What they store:

- Embedding vectors
 - Original text chunks
 - Metadata
- Role in RAG:**
- User question → embedding → similarity search → relevant chunks

Save Memory – Memory in LangChain

Problem:

- LLMs are stateless
- They forget previous conversation

Memory solves:

- Conversation history
- Context continuity
- Personalized responses

Types of Memory:

- ConversationBufferMemory
- ConversationSummaryMemory
- VectorStoreMemory

Use cases:

- Chatbots
- Assistants
- Multi-turn conversations

Retrieval Process – Generate the Answer

This is the heart of RAG

Flow:

- User asks a question
- Question converted to embedding
- Vector DB retrieves most relevant chunks
- Chunks + question passed to LLM
- LLM generates grounded answer

Benefit

- Answers are based on your data
- Reduces hallucinations
- More accurate responses

Agents Generation

What are Agents?

- Agents are decision-making LLMs.
- They can:
 - Decide which tool to use
 - Decide what step to take next
 - Perform multi-step reasoning

Example

- User asks:
 - "Analyze this PDF and summarize key points"
- Agent decides:
 - Load document
 - Chunk data
 - Retrieve relevant parts
 - Summarize

Agents = LLM + tools + reasoning loop

Chains – Connecting Everything

What are Chains?

- Chains connect: - Prompt → Model → Output **Advanced Chains:**
- RetrievalQAChain
- ConversationalRetrievalChain
- SequentialChain

Why chains are powerful:

- Automate workflows
- Connect prompts to Python functions
- Build end-to-end pipelines

Final Output and Real-World Applications

- The final output is an accurate, context-aware, and reliable answer generated using retrieved knowledge.
- Real-world applications:
 - Document Q&A systems
 - Chatbots
 - Knowledge assistants
 - Legal and medical research tools
 - Customer support automation

Conclusion

- LangChain transforms simple LLMs into powerful AI systems by integrating documents, embeddings, retrieval, memory, tools, and reasoning. -It is a foundational framework for building next-generation AI applications.

End-to-End LangChain Flow (Summary)

Document Loader ↓ Chunking ↓ Embeddings ↓ Vector Database ↓ Retriever ↓ Prompt Template ↓ LLM ↓ Memory ↓ Final Answer

In []: